

## Resnet 18 on faces

```
In [1]: import torch
import torchvision
import torchvision.transforms as transforms
import resnet
import torch.optim as optim
import torch.nn as nn
import numpy as np
import matplotlib.pyplot as plt
import random
import resnet
```

```
In [2]: batchsize = 75
rate = 0.1
epochs = 200
lr_decay = 0.85
lr_stride = 5
```

```
In [3]: class FaceDataset(torch.utils.data.Dataset):

    def __init__(self, transform, train=True):
        self.image_prefix = "face_renders/face"
        self.image_suffix = ".jpg"
        self.vertex_prefix = "processed_faces/face"
        self.vertex_suffix = ".txt"
        self.count = 5000
        self.trainn = 4500

        self.train = train
        self.transform = transform

        shape = np.loadtxt(self.vertex_prefix + str(1) + self.vertex_suffix).shape
        tmp = np.zeros((self.count, shape[0], shape[1]))
        for i in range(self.count):
            tmp[i] = np.loadtxt(self.vertex_prefix + str(i + 1) + self.vertex_suffix)

        self.mean = np.mean(tmp, axis=0)
        self.outputdim = shape[0] * shape[1]
        self.labels = [torch.from_numpy((lab - self.mean).reshape(self.outputdim)).float() for lab in tmp]

        # simple version for working with CWD

    def __len__(self):
        if self.train:
            return self.trainn
        else:
            return self.count - self.trainn

    def __getitem__(self, idx):
        if not train:
            idx += self.trainn
        y = self.labels[idx]
        x = plt.imread(self.image_prefix + str(idx + 1) + self.image_suffix)

        sample = (x,y)
        sample = (self.transform(sample[0]), sample[1])

        return sample
```

```
In [4]: transform = transforms.Compose(
        [transforms.ToTensor(),
         transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = FaceDataset(transform, train=True)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batchsize,
                                           shuffle=True, num_workers=0)

testset = FaceDataset(transform, train=False)
testloader = torch.utils.data.DataLoader(trainset, batch_size=batchsize,
                                          shuffle=True, num_workers=0)
```

```
In [5]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print("torch.cuda.is_available()  =", torch.cuda.is_available())
print("torch.cuda.device_count()  =", torch.cuda.device_count())
print("torch.cuda.device('cuda')  =", torch.cuda.device(0))
print("torch.cuda.current_device() =", torch.cuda.current_device())

def to_device(data, device):
    if isinstance(data, (list, tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)

class DeviceDataLoader():
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        return len(self.dl)

trainloader = DeviceDataLoader(trainloader, device)
testloader = DeviceDataLoader(testloader, device)

torch.cuda.is_available()  = True
torch.cuda.device_count()  = 1
torch.cuda.device('cuda')  = <torch.cuda.device object at 0x0000016199917A58>
torch.cuda.current_device() = 0
```

```
In [6]: model = resnet.resnet18(output_size=trainset.outputdim)
        model.to(device)
        optimizer = optim.SGD(model.parameters(), lr=rate)

        criterion = nn.MSELoss()

        def adjust_learning_rate(optimizer, epoch, decay, stride):
            lr = rate * (decay ** (epoch // stride))
            for param_group in optimizer.param_groups:
                param_group['lr'] = lr
```

```

In [7]: def train(model, optimizer, criterion, epochs, trainloader, testloader):
    model.train()
    samples = 1
    losses = []
    test_losses = []
    k = len(trainloader)// samples

    for epoch in range(epochs): # loop over the dataset multiple times
        running_loss = 0.0
        for i, data in enumerate(trainloader, 0):
            # get the inputs
            inputs, labels = data

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # print statistics
            running_loss += loss.item()
            if i % k == k - 1:

                losses.append(running_loss / k)

                testloss = 0
                total = 0
                iterations = 0
                with torch.no_grad():
                    for data in testloader:
                        images, labels = data
                        outputs = model(images)
                        testloss += criterion(outputs, labels)
                        total += labels.size(0)
                        iterations += 1
                        if total > 200:
                            break
                test_losses.append(testloss / iterations)

                print('[%d, %5d] loss: %.3f test_loss: %.3f' %(epoch + 1, i +
1, losses[-1], test_losses[-1]))

                running_loss = 0.0

            adjust_learning_rate(optimizer, epoch+1, lr_decay, lr_stride)

    print('Finished Training')
    plt.plot(np.arange(0, len(losses)/samples, 1.0/samples), losses)
    plt.title("loss")
    plt.xlabel("epoch")
    plt.ylabel("loss")
    plt.show()

```

```
plt.plot(np.arange(0, len(test_losses)/samples, 1.0/samples), test_losses)
plt.title("test_loss")
plt.xlabel("epoch")
plt.ylabel("test_losses")
plt.show()
model.eval()
```

```
In [8]: train(model, optimizer, criterion, epochs, trainloader, testloader)
```

```
[1, 60] loss: 9.339 test_loss: 7.488
[2, 60] loss: 7.093 test_loss: 6.447
[3, 60] loss: 5.784 test_loss: 5.349
[4, 60] loss: 5.392 test_loss: 5.143
[5, 60] loss: 5.093 test_loss: 4.929
[6, 60] loss: 4.658 test_loss: 4.444
[7, 60] loss: 4.383 test_loss: 4.296
[8, 60] loss: 4.213 test_loss: 4.098
[9, 60] loss: 4.079 test_loss: 3.975
[10, 60] loss: 3.906 test_loss: 3.811
[11, 60] loss: 3.702 test_loss: 3.654
[12, 60] loss: 3.562 test_loss: 3.593
[13, 60] loss: 3.445 test_loss: 3.304
[14, 60] loss: 3.333 test_loss: 3.271
[15, 60] loss: 3.220 test_loss: 3.213
[16, 60] loss: 3.078 test_loss: 2.941
[17, 60] loss: 2.994 test_loss: 3.029
[18, 60] loss: 2.933 test_loss: 2.827
[19, 60] loss: 2.889 test_loss: 2.789
[20, 60] loss: 2.820 test_loss: 2.685
[21, 60] loss: 2.731 test_loss: 2.643
[22, 60] loss: 2.693 test_loss: 2.666
[23, 60] loss: 2.642 test_loss: 2.540
[24, 60] loss: 2.644 test_loss: 2.582
[25, 60] loss: 2.587 test_loss: 2.599
[26, 60] loss: 2.570 test_loss: 2.493
[27, 60] loss: 2.513 test_loss: 2.439
[28, 60] loss: 2.482 test_loss: 2.446
[29, 60] loss: 2.471 test_loss: 2.388
[30, 60] loss: 2.452 test_loss: 2.439
[31, 60] loss: 2.401 test_loss: 2.450
[32, 60] loss: 2.380 test_loss: 2.346
[33, 60] loss: 2.374 test_loss: 2.327
[34, 60] loss: 2.359 test_loss: 2.298
[35, 60] loss: 2.336 test_loss: 2.390
[36, 60] loss: 2.295 test_loss: 2.210
[37, 60] loss: 2.278 test_loss: 2.227
[38, 60] loss: 2.273 test_loss: 2.222
[39, 60] loss: 2.255 test_loss: 2.332
[40, 60] loss: 2.235 test_loss: 2.204
[41, 60] loss: 2.233 test_loss: 2.109
[42, 60] loss: 2.214 test_loss: 2.166
[43, 60] loss: 2.189 test_loss: 2.126
[44, 60] loss: 2.168 test_loss: 2.177
[45, 60] loss: 2.165 test_loss: 2.170
[46, 60] loss: 2.144 test_loss: 2.210
[47, 60] loss: 2.138 test_loss: 2.140
[48, 60] loss: 2.123 test_loss: 2.066
[49, 60] loss: 2.113 test_loss: 2.029
[50, 60] loss: 2.114 test_loss: 2.107
[51, 60] loss: 2.085 test_loss: 2.036
[52, 60] loss: 2.081 test_loss: 2.090
[53, 60] loss: 2.071 test_loss: 2.140
[54, 60] loss: 2.075 test_loss: 2.208
[55, 60] loss: 2.061 test_loss: 2.091
[56, 60] loss: 2.053 test_loss: 2.006
[57, 60] loss: 2.030 test_loss: 2.012
```





```
[58, 60] loss: 2.033 test_loss: 2.026
[59, 60] loss: 2.027 test_loss: 2.071
[60, 60] loss: 2.031 test_loss: 1.967
[61, 60] loss: 2.001 test_loss: 1.956
[62, 60] loss: 2.012 test_loss: 1.955
[63, 60] loss: 2.004 test_loss: 1.931
[64, 60] loss: 1.995 test_loss: 1.933
[65, 60] loss: 2.003 test_loss: 2.020
[66, 60] loss: 1.974 test_loss: 1.986
[67, 60] loss: 1.984 test_loss: 1.984
[68, 60] loss: 1.986 test_loss: 1.893
[69, 60] loss: 1.957 test_loss: 1.992
[70, 60] loss: 1.962 test_loss: 1.902
[71, 60] loss: 1.937 test_loss: 1.984
[72, 60] loss: 1.934 test_loss: 1.977
[73, 60] loss: 1.943 test_loss: 2.013
[74, 60] loss: 1.927 test_loss: 1.898
[75, 60] loss: 1.931 test_loss: 2.019
[76, 60] loss: 1.929 test_loss: 1.980
[77, 60] loss: 1.918 test_loss: 1.910
[78, 60] loss: 1.908 test_loss: 1.953
[79, 60] loss: 1.908 test_loss: 1.889
[80, 60] loss: 1.925 test_loss: 1.911
[81, 60] loss: 1.917 test_loss: 1.893
[82, 60] loss: 1.901 test_loss: 1.975
[83, 60] loss: 1.894 test_loss: 1.895
[84, 60] loss: 1.884 test_loss: 1.887
[85, 60] loss: 1.880 test_loss: 1.940
[86, 60] loss: 1.894 test_loss: 1.930
[87, 60] loss: 1.886 test_loss: 1.848
[88, 60] loss: 1.889 test_loss: 1.914
[89, 60] loss: 1.893 test_loss: 1.919
[90, 60] loss: 1.877 test_loss: 1.888
[91, 60] loss: 1.872 test_loss: 1.938
[92, 60] loss: 1.880 test_loss: 1.880
[93, 60] loss: 1.869 test_loss: 1.888
[94, 60] loss: 1.867 test_loss: 1.809
[95, 60] loss: 1.859 test_loss: 1.882
[96, 60] loss: 1.859 test_loss: 1.823
[97, 60] loss: 1.866 test_loss: 1.815
[98, 60] loss: 1.847 test_loss: 1.784
[99, 60] loss: 1.882 test_loss: 1.845
[100, 60] loss: 1.854 test_loss: 1.833
[101, 60] loss: 1.865 test_loss: 1.834
[102, 60] loss: 1.855 test_loss: 1.854
[103, 60] loss: 1.863 test_loss: 1.842
[104, 60] loss: 1.846 test_loss: 1.894
[105, 60] loss: 1.859 test_loss: 1.875
[106, 60] loss: 1.865 test_loss: 1.857
[107, 60] loss: 1.841 test_loss: 1.835
[108, 60] loss: 1.840 test_loss: 1.777
[109, 60] loss: 1.849 test_loss: 1.816
[110, 60] loss: 1.832 test_loss: 1.779
[111, 60] loss: 1.843 test_loss: 1.781
[112, 60] loss: 1.837 test_loss: 1.871
[113, 60] loss: 1.829 test_loss: 1.765
[114, 60] loss: 1.867 test_loss: 1.797
```

```
[115, 60] loss: 1.833 test_loss: 1.919
[116, 60] loss: 1.841 test_loss: 1.761
[117, 60] loss: 1.843 test_loss: 1.829
[118, 60] loss: 1.826 test_loss: 1.721
[119, 60] loss: 1.845 test_loss: 1.811
[120, 60] loss: 1.814 test_loss: 1.819
[121, 60] loss: 1.831 test_loss: 1.800
[122, 60] loss: 1.819 test_loss: 1.804
[123, 60] loss: 1.825 test_loss: 1.795
[124, 60] loss: 1.824 test_loss: 1.955
[125, 60] loss: 1.822 test_loss: 1.920
[126, 60] loss: 1.818 test_loss: 1.834
[127, 60] loss: 1.812 test_loss: 1.817
[128, 60] loss: 1.836 test_loss: 1.799
[129, 60] loss: 1.816 test_loss: 1.866
[130, 60] loss: 1.828 test_loss: 1.801
[131, 60] loss: 1.826 test_loss: 1.845
[132, 60] loss: 1.837 test_loss: 1.779
[133, 60] loss: 1.812 test_loss: 1.777
[134, 60] loss: 1.834 test_loss: 1.857
[135, 60] loss: 1.814 test_loss: 1.771
[136, 60] loss: 1.832 test_loss: 1.747
[137, 60] loss: 1.813 test_loss: 1.729
[138, 60] loss: 1.822 test_loss: 1.809
[139, 60] loss: 1.813 test_loss: 1.832
[140, 60] loss: 1.821 test_loss: 1.785
[141, 60] loss: 1.825 test_loss: 1.816
[142, 60] loss: 1.812 test_loss: 1.827
[143, 60] loss: 1.823 test_loss: 1.797
[144, 60] loss: 1.821 test_loss: 1.777
[145, 60] loss: 1.804 test_loss: 1.826
[146, 60] loss: 1.802 test_loss: 1.775
[147, 60] loss: 1.828 test_loss: 1.823
[148, 60] loss: 1.817 test_loss: 1.800
[149, 60] loss: 1.802 test_loss: 1.768
[150, 60] loss: 1.802 test_loss: 1.789
[151, 60] loss: 1.823 test_loss: 1.759
[152, 60] loss: 1.800 test_loss: 1.813
[153, 60] loss: 1.801 test_loss: 1.771
[154, 60] loss: 1.825 test_loss: 1.788
[155, 60] loss: 1.817 test_loss: 1.763
[156, 60] loss: 1.822 test_loss: 1.768
[157, 60] loss: 1.806 test_loss: 1.921
[158, 60] loss: 1.786 test_loss: 1.734
[159, 60] loss: 1.791 test_loss: 1.810
[160, 60] loss: 1.801 test_loss: 1.864
[161, 60] loss: 1.808 test_loss: 1.842
[162, 60] loss: 1.815 test_loss: 1.819
[163, 60] loss: 1.808 test_loss: 1.882
[164, 60] loss: 1.805 test_loss: 1.803
[165, 60] loss: 1.823 test_loss: 1.719
[166, 60] loss: 1.803 test_loss: 1.734
[167, 60] loss: 1.816 test_loss: 1.779
[168, 60] loss: 1.803 test_loss: 1.907
[169, 60] loss: 1.800 test_loss: 1.791
[170, 60] loss: 1.810 test_loss: 1.783
[171, 60] loss: 1.812 test_loss: 1.826
```

```
[172, 60] loss: 1.795 test_loss: 1.793
[173, 60] loss: 1.808 test_loss: 1.803
[174, 60] loss: 1.804 test_loss: 1.828
[175, 60] loss: 1.802 test_loss: 1.862
[176, 60] loss: 1.817 test_loss: 1.823
[177, 60] loss: 1.802 test_loss: 1.813
[178, 60] loss: 1.810 test_loss: 1.874
[179, 60] loss: 1.799 test_loss: 1.770
[180, 60] loss: 1.803 test_loss: 1.870
[181, 60] loss: 1.808 test_loss: 1.755
[182, 60] loss: 1.815 test_loss: 1.705
[183, 60] loss: 1.805 test_loss: 1.803
[184, 60] loss: 1.812 test_loss: 1.924
[185, 60] loss: 1.806 test_loss: 1.791
[186, 60] loss: 1.802 test_loss: 1.755
[187, 60] loss: 1.810 test_loss: 1.796
[188, 60] loss: 1.810 test_loss: 1.803
[189, 60] loss: 1.808 test_loss: 1.783
[190, 60] loss: 1.801 test_loss: 1.762
[191, 60] loss: 1.794 test_loss: 1.836
[192, 60] loss: 1.803 test_loss: 1.796
[193, 60] loss: 1.815 test_loss: 1.729
[194, 60] loss: 1.805 test_loss: 1.793
[195, 60] loss: 1.812 test_loss: 1.768
[196, 60] loss: 1.801 test_loss: 1.751
[197, 60] loss: 1.808 test_loss: 1.763
[198, 60] loss: 1.810 test_loss: 1.807
[199, 60] loss: 1.795 test_loss: 1.880
[200, 60] loss: 1.804 test_loss: 1.801
```

Finished Training





```
In [9]: torch.save(model.state_dict(), "res18b" + str(batchsize) + "r" + str(rate) +
      "e" + str(epochs) + ".statedict")
```

```
In [11]: with torch.no_grad():
      d = next(testloader.__iter__())
      images, labels = d
      outputs = model(images)

      print(trainset.mean)
      print(outputs[0])
      print(labels[0])
      print(np.linalg.norm((labels[0] - outputs[0]).to('cpu').numpy()))
```

```
[[-57.42849171  43.50585649  81.09827847]
 [-57.10330025  40.76217867  80.77729131]
 [-56.4013172   36.99098729  79.87123442]
 ...
 [ 54.11570628 -43.99206555  66.38817782]
 [ 55.6071899  -46.20434955  60.50756982]
 [ 56.74632444 -47.9137754   53.69290953]]
```

```
tensor([-0.8509, -1.7231, -4.8252, ..., -2.9921,  2.7223,  0.9377], device
='cuda:0')
```

```
tensor([-0.5952, -1.8055, -7.2621, ..., -1.7967,  2.6882,  0.4817], device
='cuda:0')
```

```
56.56403
```