

Resnet 18 on faces

```
In [1]: import torch
import torchvision
import torchvision.transforms as transforms
import torch.optim as optim
import torch.nn as nn
import numpy as np
import matplotlib.pyplot as plt
import random
import resnet
```

```
In [11]: batchsize = 75
rate = 0.15
epochs = 150
lr_decay = 0.84
lr_stride = 5
```

```
In [3]: class FaceDataset(torch.utils.data.Dataset):

    def __init__(self, transform, train=True):
        self.image_prefix = "face_renderers/face"
        self.image_suffix = ".jpg"
        self.vertex_prefix = "processed_faces/face"
        self.vertex_suffix = ".txt"
        self.count = 5000
        self.trainn = 4500

        self.train = train
        self.transform = transform

        shape = np.loadtxt(self.vertex_prefix + str(1) + self.vertex_suffix).shape
        tmp = np.zeros((self.count, shape[0], shape[1]))
        for i in range(self.count):
            tmp[i] = np.loadtxt(self.vertex_prefix + str(i + 1) + self.vertex_suffix)

        self.mean = np.mean(tmp, axis=0)
        self.outputdim = shape[0] * shape[1]
        self.labels = [torch.from_numpy((lab - self.mean).reshape(self.outputdim)).float() for lab in tmp]

        self.images = [plt.imread(self.image_prefix + str(i + 1) + self.image_suffix) for i in range(self.count)]

        # simple version for working with CWD

    def __len__(self):
        if self.train:
            return self.trainn
        else:
            return self.count - self.trainn

    def __getitem__(self, idx):
        if not self.train:
            idx += self.trainn
        return (self.transform(self.images[idx]), self.labels[idx])
```

```
In [4]: transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = FaceDataset(transform, train=True)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batchsize,
                                           shuffle=True, num_workers=0)

testset = FaceDataset(transform, train=False)
testloader = torch.utils.data.DataLoader(testset, batch_size=batchsize,
                                          shuffle=False, num_workers=0)
```

```
In [5]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print("torch.cuda.is_available()  =", torch.cuda.is_available())
print("torch.cuda.device_count()  =", torch.cuda.device_count())
print("torch.cuda.device('cuda')  =", torch.cuda.device(0))
print("torch.cuda.current_device() =", torch.cuda.current_device())
```

```
def to_device(data, device):
    if isinstance(data, (list, tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)
```

```
class DeviceDataLoader():
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        return len(self.dl)
```

```
trainloader = DeviceDataLoader(trainloader, device)
testloader = DeviceDataLoader(testloader, device)
```

```
torch.cuda.is_available()  = True
torch.cuda.device_count()  = 1
torch.cuda.device('cuda')  = <torch.cuda.device object at 0x0000019D7A45EE10>
torch.cuda.current_device() = 0
```

```
In [18]: model = resnet.resnet18(output_size=trainset.outputdim)
model.to(device)
#optimizer = optim.SGD(model.parameters(), lr=rate)
#optimizer = optim.Adam(model.parameters(), lr=rate)
optimizer = optim.SGD(model.parameters(), lr=rate, momentum=0.95)

criterion = nn.MSELoss()

def adjust_learning_rate(optimizer, epoch, decay, stride):
    lr = rate * (decay ** (epoch // stride))
    for param_group in optimizer.param_groups:
        param_group['lr'] = lr
```

```

In [19]: def train(model, optimizer, criterion, epochs, trainloader, testloader):
    model.train()
    samples = 1
    losses = []
    test_losses = []
    k = len(trainloader)// samples

    for epoch in range(epochs): # loop over the dataset multiple times
        running_loss = 0.0
        for i, data in enumerate(trainloader, 0):
            # get the inputs
            inputs, labels = data

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # print statistics
            running_loss += loss.item()
            if i % k == k - 1:

                losses.append(running_loss / k)

                testloss = 0
                total = 0
                iterations = 0
                with torch.no_grad():
                    for data in testloader:
                        images, labels = data
                        outputs = model(images)
                        testloss += criterion(outputs, labels)
                        total += labels.size(0)
                        iterations += 1
                        if total > 200:
                            break
                test_losses.append(testloss / iterations)

                print('[%d, %5d] loss: %.3f test_loss: %.3f' %(epoch + 1, i +
1, losses[-1], test_losses[-1]))

                running_loss = 0.0

            adjust_learning_rate(optimizer, epoch+1, lr_decay, lr_stride)

    print('Finished Training')
    plt.plot(np.arange(0, len(losses)/samples, 1.0/samples), losses)
    plt.title("loss")
    plt.xlabel("epoch")
    plt.ylabel("loss")
    plt.show()

```

```
plt.plot(np.arange(0, len(test_losses)/samples, 1.0/samples), test_losses)
plt.title("test_loss")
plt.xlabel("epoch")
plt.ylabel("test_losses")
plt.show()
model.eval()
```

```
In [20]: train(model, optimizer, criterion, epochs, trainloader, testloader)
```


```
[1, 60] loss: 7.306 test_loss: 5.625
[2, 60] loss: 5.029 test_loss: 4.915
[3, 60] loss: 4.577 test_loss: 4.466
[4, 60] loss: 4.170 test_loss: 4.125
[5, 60] loss: 3.817 test_loss: 3.933
[6, 60] loss: 3.526 test_loss: 3.733
[7, 60] loss: 3.362 test_loss: 3.611
[8, 60] loss: 3.235 test_loss: 3.522
[9, 60] loss: 3.123 test_loss: 3.416
[10, 60] loss: 3.076 test_loss: 3.353
[11, 60] loss: 2.882 test_loss: 3.217
[12, 60] loss: 2.762 test_loss: 3.201
[13, 60] loss: 2.713 test_loss: 3.248
[14, 60] loss: 2.673 test_loss: 3.150
[15, 60] loss: 2.566 test_loss: 3.113
[16, 60] loss: 2.495 test_loss: 3.103
[17, 60] loss: 2.355 test_loss: 2.961
[18, 60] loss: 2.260 test_loss: 2.934
[19, 60] loss: 2.215 test_loss: 2.981
[20, 60] loss: 2.139 test_loss: 2.991
[21, 60] loss: 2.064 test_loss: 2.962
[22, 60] loss: 1.977 test_loss: 2.791
[23, 60] loss: 1.897 test_loss: 2.838
[24, 60] loss: 1.826 test_loss: 2.707
[25, 60] loss: 1.792 test_loss: 2.729
[26, 60] loss: 1.682 test_loss: 2.653
[27, 60] loss: 1.598 test_loss: 2.652
[28, 60] loss: 1.527 test_loss: 2.649
[29, 60] loss: 1.473 test_loss: 2.633
[30, 60] loss: 1.433 test_loss: 2.686
[31, 60] loss: 1.346 test_loss: 2.634
[32, 60] loss: 1.294 test_loss: 2.579
[33, 60] loss: 1.224 test_loss: 2.609
[34, 60] loss: 1.206 test_loss: 2.608
[35, 60] loss: 1.164 test_loss: 2.571
[36, 60] loss: 1.110 test_loss: 2.540
[37, 60] loss: 1.067 test_loss: 2.572
[38, 60] loss: 1.054 test_loss: 2.549
[39, 60] loss: 1.019 test_loss: 2.514
[40, 60] loss: 0.998 test_loss: 2.526
[41, 60] loss: 0.947 test_loss: 2.504
[42, 60] loss: 0.960 test_loss: 2.550
[43, 60] loss: 0.913 test_loss: 2.507
[44, 60] loss: 0.903 test_loss: 2.510
[45, 60] loss: 0.863 test_loss: 2.497
[46, 60] loss: 0.849 test_loss: 2.486
[47, 60] loss: 0.829 test_loss: 2.468
[48, 60] loss: 0.810 test_loss: 2.450
[49, 60] loss: 0.791 test_loss: 2.464
[50, 60] loss: 0.770 test_loss: 2.447
[51, 60] loss: 0.753 test_loss: 2.442
[52, 60] loss: 0.738 test_loss: 2.458
[53, 60] loss: 0.730 test_loss: 2.462
[54, 60] loss: 0.721 test_loss: 2.442
[55, 60] loss: 0.721 test_loss: 2.456
[56, 60] loss: 0.698 test_loss: 2.455
[57, 60] loss: 0.677 test_loss: 2.437
```


```
[58, 60] loss: 0.684 test_loss: 2.490
[59, 60] loss: 0.668 test_loss: 2.429
[60, 60] loss: 0.673 test_loss: 2.465
[61, 60] loss: 0.653 test_loss: 2.459
[62, 60] loss: 0.639 test_loss: 2.446
[63, 60] loss: 0.639 test_loss: 2.442
[64, 60] loss: 0.657 test_loss: 2.455
[65, 60] loss: 0.628 test_loss: 2.440
[66, 60] loss: 0.603 test_loss: 2.434
[67, 60] loss: 0.598 test_loss: 2.437
[68, 60] loss: 0.599 test_loss: 2.442
[69, 60] loss: 0.587 test_loss: 2.440
[70, 60] loss: 0.583 test_loss: 2.433
[71, 60] loss: 0.581 test_loss: 2.433
[72, 60] loss: 0.573 test_loss: 2.438
[73, 60] loss: 0.571 test_loss: 2.425
[74, 60] loss: 0.563 test_loss: 2.436
[75, 60] loss: 0.563 test_loss: 2.444
[76, 60] loss: 0.555 test_loss: 2.434
[77, 60] loss: 0.552 test_loss: 2.426
[78, 60] loss: 0.555 test_loss: 2.434
[79, 60] loss: 0.550 test_loss: 2.448
[80, 60] loss: 0.552 test_loss: 2.433
[81, 60] loss: 0.539 test_loss: 2.432
[82, 60] loss: 0.529 test_loss: 2.432
[83, 60] loss: 0.532 test_loss: 2.437
[84, 60] loss: 0.531 test_loss: 2.450
[85, 60] loss: 0.523 test_loss: 2.436
[86, 60] loss: 0.532 test_loss: 2.431
[87, 60] loss: 0.525 test_loss: 2.442
[88, 60] loss: 0.524 test_loss: 2.430
[89, 60] loss: 0.514 test_loss: 2.436
[90, 60] loss: 0.516 test_loss: 2.443
[91, 60] loss: 0.510 test_loss: 2.431
[92, 60] loss: 0.504 test_loss: 2.431
[93, 60] loss: 0.507 test_loss: 2.434
[94, 60] loss: 0.508 test_loss: 2.435
[95, 60] loss: 0.497 test_loss: 2.436
[96, 60] loss: 0.497 test_loss: 2.430
[97, 60] loss: 0.491 test_loss: 2.431
[98, 60] loss: 0.492 test_loss: 2.438
[99, 60] loss: 0.492 test_loss: 2.439
[100, 60] loss: 0.495 test_loss: 2.439
[101, 60] loss: 0.483 test_loss: 2.435
[102, 60] loss: 0.489 test_loss: 2.440
[103, 60] loss: 0.482 test_loss: 2.441
[104, 60] loss: 0.485 test_loss: 2.428
[105, 60] loss: 0.491 test_loss: 2.439
[106, 60] loss: 0.495 test_loss: 2.437
[107, 60] loss: 0.484 test_loss: 2.430
[108, 60] loss: 0.479 test_loss: 2.433
[109, 60] loss: 0.481 test_loss: 2.437
[110, 60] loss: 0.479 test_loss: 2.437
[111, 60] loss: 0.478 test_loss: 2.431
[112, 60] loss: 0.472 test_loss: 2.437
[113, 60] loss: 0.475 test_loss: 2.435
[114, 60] loss: 0.472 test_loss: 2.434
```



```
[115, 60] loss: 0.469 test_loss: 2.436
[116, 60] loss: 0.468 test_loss: 2.439
[117, 60] loss: 0.474 test_loss: 2.437
[118, 60] loss: 0.470 test_loss: 2.434
[119, 60] loss: 0.466 test_loss: 2.441
[120, 60] loss: 0.467 test_loss: 2.440
[121, 60] loss: 0.474 test_loss: 2.440
[122, 60] loss: 0.465 test_loss: 2.438
[123, 60] loss: 0.470 test_loss: 2.439
[124, 60] loss: 0.469 test_loss: 2.444
[125, 60] loss: 0.469 test_loss: 2.437
[126, 60] loss: 0.468 test_loss: 2.439
[127, 60] loss: 0.462 test_loss: 2.438
[128, 60] loss: 0.462 test_loss: 2.439
[129, 60] loss: 0.465 test_loss: 2.436
[130, 60] loss: 0.461 test_loss: 2.437
[131, 60] loss: 0.461 test_loss: 2.437
[132, 60] loss: 0.459 test_loss: 2.439
[133, 60] loss: 0.457 test_loss: 2.441
[134, 60] loss: 0.457 test_loss: 2.442
[135, 60] loss: 0.452 test_loss: 2.440
[136, 60] loss: 0.457 test_loss: 2.441
[137, 60] loss: 0.460 test_loss: 2.439
[138, 60] loss: 0.459 test_loss: 2.441
[139, 60] loss: 0.463 test_loss: 2.439
[140, 60] loss: 0.455 test_loss: 2.440
[141, 60] loss: 0.463 test_loss: 2.440
[142, 60] loss: 0.454 test_loss: 2.440
[143, 60] loss: 0.454 test_loss: 2.439
[144, 60] loss: 0.452 test_loss: 2.440
[145, 60] loss: 0.459 test_loss: 2.440
[146, 60] loss: 0.454 test_loss: 2.441
[147, 60] loss: 0.452 test_loss: 2.440
[148, 60] loss: 0.452 test_loss: 2.440
[149, 60] loss: 0.459 test_loss: 2.441
[150, 60] loss: 0.453 test_loss: 2.439
```

Finished Training





```
In [21]: torch.save(model.state_dict(), "res18b" + str(batchsize) + "r" + str(rate) +  
"e" + str(epochs) + ".statedict")
```

```
In [ ]: with torch.no_grad():  
    d = next(testloader.__iter__())  
    images, labels = d  
    outputs = model(images)  
  
    print(trainset.mean)  
    print(outputs[0])  
    print(labels[0])  
    print(np.linalg.norm((labels[0] - outputs[0]).to('cpu').numpy()))
```

```
In [ ]: np.mean(np.std([l.numpy() for l in trainset.labels], axis=0))
```

```
In [ ]: model = resnet.resnet101(output_size=trainset.outputdim)  
model.load_state_dict(torch.load("res101b75r0.15e150.statedict"))  
# model.to(device)  
model.eval()
```

```
In [26]: running_loss = 0
criterion = nn.MSELoss()
start = 4501
for i, data in enumerate(testloader, 0):
    # get the inputs
    inputs, labels = data

    # forward + backward + optimize
    outputs = model(inputs)

    for o in outputs:
        arr = o.detach().to("cpu").numpy().reshape((564,3)) + testset.mean
        np.savetxt("res50predictions/face" + str(start) + ".txt", arr, fmt="%.
9f")
        start += 1

    loss = criterion(outputs, labels)

    # print statistics
    running_loss += loss.item() * len(outputs)

print(running_loss / 500 * trainset.outputdim)

4014.2447912693024
```

```
In [ ]: print (len(testloader))
```

```
In [ ]: output = model(testset.transform(plt.imread("real/IMG_3368_cropped.png")[:, :, :
3])).unsqueeze(0))
output = output.detach().to("cpu").numpy().reshape((564,3)) + testset.mean
np.savetxt("res101predictions/facecustom.txt", output, fmt="%.9f")
```